

Design and Implementation of a Intelligent Virtual Agent for Kabré Supercomputer Infrastructure Support

Dylan Benavides^{1,2}[0009–0006–8777–7876]

¹ National High Technology Center, San José 10109, Costa Rica
dbenavides@cenat.ac.cr

<https://www.cenat.ac.cr/es/>

² University of Costa Rica, San José 11501-2060, Costa Rica
dylan.benavides@ucr.ac.cr

<https://www.ucr.ac.cr>

Abstract. High-Performance Computing (HPC) environments, such as the Kabré supercomputer at the National High Technology Center (CENAT) in Costa Rica, face significant challenges in providing efficient and scalable user support due to the heterogeneity and growing size of their user communities. Traditional ticket-based systems introduce delays, repetitive workload for support staff, and limited interactivity, negatively impacting both user experience and system efficiency.

This work proposes the design, implementation, and evaluation of an Intelligent Virtual Agent (IVA) based on large language models (LLMs), integrated into the Kabré infrastructure to provide real-time technical assistance. The agent is designed to handle common support tasks such as job execution issues, software configuration, and resource management, while escalating complex cases to human experts when necessary.

The system is built on a multi-agent architecture that leverages natural language processing (NLP) techniques applied to the historical ticket database of the Kabré support system. This enables intent classification and dynamic routing of user requests to specialized agents according to the detected problem category.

User interaction is performed exclusively through text input, while system responses are delivered in both textual form and through text-to-speech (TTS) synthesis, providing a multimodal output experience without requiring voice input from the user.

Finally, the system was evaluated through a within-subject experimental study using representative HPC support tasks, comparing the performance of the IVA against users' prior experience with the traditional ticketing system. The evaluation was conducted in terms of efficiency, usability, and cognitive load using objective performance metrics and standardized questionnaires.

Keywords: High-Performance Computing (HPC) · Intelligent Virtual Agent · Large Language Models · Kabré supercomputer · technical support systems · human-computer interaction · conversational agents · user experience.

1 Problem Definition

High-Performance Computing (HPC) environments typically host highly heterogeneous user communities composed of students, researchers, and professionals from diverse disciplines. This diversity naturally leads to a wide range of technical expertise, resulting in queries of varying complexity—from basic questions about system usage and command-line operations to advanced issues involving software installation, resource optimization, and memory management.

In Costa Rica, the Kabré supercomputer, operated by the National High Technology Center (CeNAT), represents a critical infrastructure for scientific research. The system currently supports approximately 300 active users, generating over 100,000 computing hours annually. This community is primarily composed of students, faculty members, and researchers from public universities across the country, spanning multiple domains of knowledge. Despite its importance, user support poses a significant operational challenge, as it is managed by a small team consisting of a single infrastructure engineer and one technical assistant.

At present, user support is handled through a ticket-based system, where users submit queries that may include problem descriptions, code snippets, and supporting files. While functional, this approach presents several limitations. First, it introduces friction in the user experience, as users must navigate the platform, authenticate, and formally structure their requests before receiving assistance. Second, resolving issues often requires multiple back-and-forth interactions to clarify context, which increases response and resolution times.

Furthermore, the system becomes overloaded due to both the volume of requests and the recurrence of similar queries. Many of these correspond to frequently asked questions that require repetitive responses, leading to inefficient use of the support team’s time. At the same time, more specialized queries—such as those arising in fields like bioinformatics or computational sciences—require in-depth, domain-specific analysis. This combination of routine and complex queries contributes to system inefficiency and the accumulation of unresolved tickets.

Although a knowledge base exists within the current system, it is outdated, limited in scope, and lacks interactivity. Users cannot perform targeted queries or obtain immediate answers, forcing them to rely entirely on the ticketing system. As a result, response times can exceed 24 hours, which is particularly problematic in scenarios requiring rapid troubleshooting or near real-time assistance.

Given these challenges, there is a clear need to leverage modern technological solutions to improve user support in HPC environments. In particular, the development of an intelligent virtual agent emerges as a promising approach to automate the handling of user queries, especially those that are repetitive in nature. Such a system could provide real-time assistance, reduce the operational burden on infrastructure staff, and enhance the overall user experience. However, in its initial stages, the agent should not operate with full autonomy; instead, it should work in conjunction with human experts, determining which queries can

be effectively resolved automatically and which require escalation to the support team.

2 Literature Review

Early research in High-Performance Computing (HPC) began to challenge the notion that computational performance alone defines system effectiveness, instead emphasizing the importance of interaction, usability, and human-centered design. One of the earliest contributions in this direction is the work of Faigle, Fox, Furmanski, Niemiec, and Simoni (1993), who established that virtual reality (VR) and HPC are inherently complementary technologies. They argued that large-scale “Grand Challenge” simulations require dynamic VR interfaces capable of visualizing complex computational processes in real time. A key technical insight from their work is the need to decouple the update rates of HPC computational engines from VR display systems to prevent instability and user discomfort, such as simulation sickness [9]. Despite these advances, the authors identified critical limitations, including the lack of standardized connectivity models and the absence of hardware-independent interface protocols. These challenges remain relevant today, highlighting a persistent gap between computational performance and interactive usability that modern HPC systems continue to address [9].

Building on the importance of interaction, Cai and He (2004) introduced an early approach to enhancing human-machine communication through avatar-based interfaces. At the time, most interaction systems relied on rigid rule-based mechanisms and static learning models that were unable to adapt to different users or contexts. To overcome these limitations, the authors proposed the Interactive Agent Based on Awareness (ABI Agent), a system designed to simulate human-like behavior through a structured process of intention generation, goal selection, and action planning. Their results demonstrated that embedding behavioral characteristics into agents can significantly improve user efficiency by automating repetitive tasks and enabling more intuitive interaction.

However, Cai and He (2004) also acknowledged fundamental challenges in achieving truly adaptive and human-like systems. As discussed in [4], integrating personality and human knowledge into intelligent agents remains difficult due to the complexity and variability of psychological and contextual factors across users. Additionally, early systems lacked the ability to learn autonomously from interaction, limiting their capacity to adapt dynamically to user needs. These limitations underscore the broader challenge of moving from static, rule-based systems toward more flexible, learning-driven interaction models.

A more explicit focus on the human dimension in HPC emerges in the work of Monroy, Becerra, Bellas, and Duro (2006, 2008), who argue that traditional optimization strategies are often misaligned with user needs. Rather than prioritizing metrics such as job turnaround time or queue waiting time, they emphasize that the ultimate objective of HPC systems should be to maximize user satisfaction while minimizing inefficiencies. Their studies reveal that users frequently

overestimate required computational resources—particularly CPU time—as a precaution against job failure, leading to suboptimal scheduling and reduced system efficiency.

To address this issue, the authors developed the EVOPROC interface, which combines neural networks and evolutionary algorithms to predict actual resource consumption more accurately than manual user submissions [15, 14]. This approach enables HPC systems to interpret imprecise user inputs and translate them into optimized execution plans, effectively bridging the gap between human intent and system-level performance. Furthermore, they identify behavioral indicators such as frequent system queries, including `qstat`, as signals of user uncertainty and impatience, highlighting the importance of incorporating behavioral data into system design.

Subsequent research, including [16], reinforces the importance of user behavior modeling as a critical yet underdeveloped component of HPC scheduling systems. In particular, these studies point out the difficulty of formalizing subjective dimensions such as satisfaction, trust, and perceived responsiveness, which vary significantly across users and contexts. As a result, a persistent gap remains between technical performance metrics and human-centered evaluation.

Moreover, both Monroy et al. and later studies highlight a structural limitation in traditional HPC schedulers: although they provide extensive configuration parameters, there is often no clear mapping between these parameters and the quality of service perceived by users. This disconnect is further exacerbated by users’ limited understanding of their own computational requirements. Consequently, there is a growing need for intelligent systems capable of translating uncertain or conservative user inputs into efficient and reliable execution strategies.

Overall, these foundational works mark a critical transition in HPC research—from systems focused purely on computational optimization toward environments that increasingly account for human interaction, perception, and adaptability. While early approaches laid the groundwork for integrating usability into HPC, they also revealed significant limitations that continue to motivate the development of more intelligent, user-aware, and autonomous systems.

To bridge the early human-centered vision of HPC with more mature intelligent interaction systems, recent research begins to formalize how user experience, personalization, and adaptive assistance can be operationalized at scale. While foundational works established the importance of visualization, user satisfaction, and basic agent behavior, the next stage of research shifts toward embedding intelligence directly into support systems, enabling them to interpret user needs, learn from interaction data, and provide context-aware assistance. This transition marks the emergence of a second layer in HPC evolution, where interaction is no longer passive but actively mediated by intelligent, adaptive systems.

Within this context, Biggs and Dalton (2024) highlight a persistent knowledge gap among researchers, emphasizing that expertise in one HPC system does not easily transfer to another. To address this, they demonstrate that integrating AI assistants based on *RAG* enables highly precise, site-specific technical sup-

port grounded in continuously updated internal documentation, thus avoiding the costly need for frequent model retraining [2]. However, they also identify a critical limitation: historical support data—such as tickets—often contain noise, including irrelevant conversations or sensitive information, which complicates their transformation into reliable knowledge sources.

This challenge of creating more natural and adaptive interaction systems echoes earlier concerns raised by Cai and He (2004), who argue that traditional avatar-based interfaces lack personality and flexibility. Their proposed ABI Agent represents an important step toward more human-like interaction, introducing a structured process of intention generation, goal selection, and action planning. By embedding behavioral tendencies into agents, their approach demonstrates measurable improvements in user efficiency, particularly in handling repetitive or cognitively demanding tasks. Nevertheless, they acknowledge that fully modeling human personality and contextual understanding remains a complex and unresolved problem, especially given the variability of psychological and sensory factors across users [4].

Expanding on the limitations of interaction design, Diederich et al. (2020) provide a broader perspective on conversational agents, arguing that many systems fail not because of technical shortcomings, but due to their neglect of user-specific characteristics such as culture, personality, and expectations. They advocate for longitudinal, real-world studies to better understand how social and behavioral cues influence long-term user engagement and system sustainability [6]. This reinforces the idea that effective HPC support systems must go beyond functional correctness and incorporate deeper human-centered design principles.

From an operational standpoint, Sanjeewa (2020) situates these challenges within the broader landscape of technical support systems, noting that the growing complexity of enterprise environments places increasing pressure on support infrastructures. The study emphasizes that AI-driven support tools are most effective when deployed in a human-in-the-loop configuration, where chatbots handle repetitive queries while experts focus on complex issues. However, important gaps remain, including the lack of specialized ticketing systems and the need for continuous learning mechanisms driven by real user feedback, as well as enhanced security features such as multi-factor authentication.

Building on these ideas within the HPC domain, Shifare, Shukla, and Kothari (2025) demonstrate how intelligent systems can significantly improve both access management and user support. Their AI-based Laboratory Manager (ALM) reduces resource access approval times from days to minutes, while a fine-tuned RoBERTa-based chatbot achieves a 98.1% accuracy rate in resolving user queries. Despite these advances, the study highlights persistent issues such as ambiguity in user requests and inconsistencies in documentation, while also pointing toward future directions in predictive resource management and more advanced scheduling optimization [20].

Finally, Wulf and Meierhofer (2024) extend this discussion by analyzing the role of large language models in technical customer service environments. They show that while LLMs are highly effective for low-level cognitive tasks such

as summarization and translation, more complex problem-solving requires enhanced architectures like *RAG* or fine-tuning to reduce hallucinations. Their work also underscores the importance of structured data ecosystems to enable context-aware AI systems. However, they identify a critical gap in the lack of large-scale, real-world validation studies, as well as challenges related to user acceptance, usability, and integration with legacy systems [21].

Together, these studies define Cluster 2 as a transition toward intelligent, user-centered support systems in HPC, where the focus shifts from static interfaces to adaptive, learning-driven interaction layers capable of bridging the gap between complex computational infrastructures and diverse user communities.

As intelligent interaction layers mature and begin to effectively mediate between users and complex HPC infrastructures, the next logical step is to optimize what happens behind the interface: the allocation, scheduling, and orchestration of computational resources. This shift marks a transition from user-centered assistance toward system-level intelligence, where AI is not only supporting users but actively managing workloads, predicting behavior, and optimizing performance in real time. In this sense, Cluster 3 establishes the operational backbone upon which higher-level automation—such as fully autonomous support systems—can be built.

Within this framework, Fan et al. (2020) demonstrate that traditional scheduling heuristics like FCFS are fundamentally inadequate for dynamic HPC environments, as they fail to adapt to fluctuating workloads and often require continuous manual tuning. Their proposed DRAS system leverages hierarchical deep reinforcement learning to autonomously manage resource reservation and back-filling, achieving performance improvements of up to 45%. However, the authors highlight a critical limitation: models originally designed for cloud computing do not translate well to HPC due to factors such as job rigidity and starvation, revealing an important gap in domain-specific AI design [10].

Complementing this work, Narantuya et al. (2022) expand the optimization perspective by incorporating energy efficiency into the scheduling problem. Their decentralized multi-agent deep reinforcement learning (*mDRL*) approach significantly improves both performance and sustainability, reducing task completion times by 20% and inefficient energy consumption by 40% [17]. Despite these gains, the study exposes structural limitations in current HPC ecosystems, particularly regarding accessibility for smaller institutions and the inability of centralized architectures to meet low-latency demands. This reinforces the need for hybrid models that integrate Edge Computing with traditional HPC infrastructures.

From a behavioral standpoint, Zaborovsky et al. (2023) reveal that inefficiencies in HPC systems are not only technical but also human-driven. Their analysis shows that users overestimate job execution times in more than 89% of cases, leading to poor resource utilization and congested queues. To address this, they propose a multi-agent scheduling system enhanced with Random Survival Forests (RSF), which predicts probabilistic execution times rather than fixed estimates. While this improves accuracy, the authors emphasize the neces-

sity of Explainable AI (XAI) to ensure that users understand system decisions, and they advocate for the integration of Natural Language Processing (NLP) interfaces to facilitate more intuitive interaction with scheduling systems [23].

In parallel, Jain, Gupta, and Neha (2024) tackle inefficiencies from the perspective of support infrastructure, arguing that traditional ticket management systems are ill-equipped to handle the complexity of modern IT environments. Their multidimensional clustering approach—combining semantic, spatial, and temporal features—achieves a fidelity of 96% while reducing manual workload by 30% [12]. However, they identify a major barrier in the transition from legacy systems to AI-driven frameworks, particularly in preserving institutional knowledge and overcoming organizational resistance.

At the architectural level, Oparin, Bogdanova, and Pashinin (2020) introduce the *HPC SOMAS* framework, which integrates hierarchical control with decentralized peer-to-peer agent networks. This hybrid approach enables high scalability and flexibility in solving complex distributed problems, yet it also highlights the growing complexity of such systems, especially for non-expert users. The lack of high-level abstraction tools remains a key obstacle for broader adoption [18].

Finally, Rosendo et al. (2025) push the boundaries of automation by proposing AI-driven orchestration of scientific workflows through modular agent-based architectures. By combining natural language interfaces with programmable APIs, their system enables adaptive and near real-time coordination of computational resources, significantly accelerating scientific experimentation [19]. Nonetheless, they identify unresolved challenges related to cross-facility integration, as well as limitations imposed by existing security policies that are not designed for interactive or streaming workflows.

Together, these contributions define Cluster 3 as the core of AI-driven optimization in HPC, where scheduling, resource management, and workflow orchestration converge into intelligent, adaptive systems.

Building upon this operational intelligence, the next stage of evolution—Cluster 4—shifts the focus toward closing the loop between system optimization and user support. While Cluster 3 optimizes how resources are allocated and managed, Cluster 4 explores how large language models, retrieval systems, and conversational agents can make these complex systems accessible, interpretable, and increasingly autonomous from the user’s perspective.

Building upon the advances in AI-driven scheduling and resource optimization, the final stage in this evolution focuses on closing the loop between system intelligence and user interaction. While Cluster 3 demonstrates how HPC infrastructures can autonomously manage workloads and optimize performance, Cluster 4 shifts attention toward making these capabilities accessible, interpretable, and actionable for users. In this layer, large language models (LLMs), retrieval-augmented generation (RAG), and conversational agents emerge as the key enablers of fully integrated, intelligent support ecosystems—where interaction, automation, and decision-making converge into a unified experience.

At the core of this transition, Rosendo et al. (2025) propose a modular, agent-based architecture that bridges workflow orchestration with natural language interfaces and programmable APIs. Their approach enables adaptive and explainable coordination of scientific workflows in near real time, effectively reducing the reliance on manual intervention [19]. However, they identify persistent barriers in cross-facility integration, particularly due to misaligned security policies and the lack of standards for interactive or streaming workloads—highlighting a structural limitation in current HPC ecosystems.

Complementing this systems-level perspective, Wulf and Meierhofer (2024) explore the role of LLMs in automating technical support tasks. They demonstrate that while models such as GPT-4 are highly effective for low-level cognitive operations (e.g., summarization and translation), more complex reasoning tasks require enhanced architectures like RAG or fine-tuning to mitigate hallucinations. Their work underscores the importance of structured data ecosystems for enabling context-aware AI, yet also reveals a critical gap in large-scale, real-world validation and system integration [21].

Addressing the challenge of fragmented knowledge, Bondapalli et al. (2025) show that HPC documentation—often distributed across heterogeneous formats—can be effectively leveraged through multimodal RAG systems. By preserving the structural integrity of technical content such as code and tables, their approach significantly improves response relevance, achieving a score of 4.7/5.0 [7]. Nevertheless, they identify that retrieval quality remains highly sensitive to noisy or irrelevant context, motivating the need for more advanced techniques such as hybrid search and knowledge graph integration.

From an applied support perspective, Pasupuleti et al. (2025) demonstrate the practical impact of LLM-based chatbots in HPC environments. Their system, enhanced with FAISS and BERT-based semantic retrieval, is capable of autonomously resolving up to 70

Similarly, Shifare et al. (2025) present an integrated solution combining an AI-based Laboratory Manager (ALM) with a RoBERTa-based chatbot, achieving a 98.1

Extending beyond support automation, Yin et al. (2024) introduce the chatHPC architecture, demonstrating how LLMs can be fully integrated into HPC infrastructures to enable programmable and conversational scientific workflows. Their results show significant improvements in model training efficiency and response accuracy through a hybrid combination of supervised fine-tuning and RAG. However, they point out unresolved challenges in scalable model optimization, dataset generation, and hallucination mitigation, as well as the need for tighter integration with simulation frameworks for end-to-end automation [22].

From a communication standpoint, Guo et al. (2021) contribute to the development of more natural interaction modalities by enhancing text-to-speech systems with contextual and semantic encoders. Their approach enables more human-like conversational agents capable of capturing spontaneous speech patterns, which is essential for immersive and effective user support [3]. Complemen-

tarily, Ashish et al. (2025) expand the role of AI agents into the domain of human collaboration, demonstrating how virtual mentors powered by lightweight LLMs can facilitate intercultural communication in global HPC projects, although challenges remain in long-term evaluation and curricular integration [13].

Finally, Jain et al. (2024) reinforce the importance of intelligent support infrastructures by showing how multidimensional AI models can significantly improve incident management through more accurate ticket clustering and resolution recommendations. Despite achieving high performance, their work highlights ongoing challenges in adaptability, contextual reasoning, and the transition from legacy systems—issues that directly impact the deployment of LLM-based support ecosystems at scale [8].

Taken together, these contributions define Cluster 4 as the most advanced and integrative layer in the evolution of HPC systems. Here, intelligent agents, LLMs, and retrieval-based architectures converge to create adaptive, context-aware, and increasingly autonomous support ecosystems. Rather than simply assisting users or optimizing resources in isolation, these systems aim to unify interaction, computation, and knowledge into a cohesive framework capable of supporting the full lifecycle of scientific discovery.

3 Research Questions

Based on the challenges identified in the Kabré HPC support system and the structural requirements of Human-Computer Interaction (HCI) methodologies, this study explores the design, implementation, and user experience of an Intelligent Virtual Agent (IVA). The research is explicitly guided by the following core research questions:

1. To what extent can an Intelligent Virtual Agent improve the efficiency of technical support in the Kabré supercomputer environment, specifically minimizing user-side waiting latencies and optimizing request processing compared to traditional ticketing?
2. What is the perceived usability of an automated IVA across a user base with diverse levels of technical expertise, and how does the agent’s operational autonomy influence user trust when navigating complex supercomputing tasks?
3. Does interacting with a 3D animated avatar that speaks and shows expressions actually help reduce user frustration, or does it feel like an unnecessary distraction compared to a simple, text-only chat?

Beyond the core user experience assessment, the study addresses critical operational and structural dimensions of the system. First, it examines how autonomy can be safely architected so the agent can accurately decide which technical queries to resolve automatically and which to escalate to CeNAT human experts. This operational boundary is heavily tied to security considerations, particularly in defining safeguards and sandboxes for scenarios where the IVA interacts with cluster configurations, ensuring zero risk to the HPC infrastructure. Finally, the study evaluates the underlying data pipeline, investigating how data quality and

curation from domain experts influence both the language model’s technical reliability and its capacity to maintain a clear, supportive communication style across a diverse user base.

4 Objectives of the Agent

The Intelligent Virtual Agent (IVA) is designed to act as a first-line support system for the Kabré High-Performance Computing (HPC) infrastructure. Its main objectives are:

1. To provide real-time assistance to users by resolving common technical issues related to system usage, job execution, software installation, and resource management.
2. To reduce the workload of the human support team by handling repetitive and frequently asked questions, while ensuring that complex cases are escalated when necessary.
3. To improve the overall user experience within the HPC environment by offering clear, adaptive, and accessible guidance tailored to different user profiles.

5 Resource Management and Coexistence in HPC Computing

agregar aquí donde se instalarán los modelos nodos etc y el protocolo coordinar con Herson

6 Preliminary Technical Architecture

The architecture of the Intelligent Virtual Agent (IVA) is designed as an integrated system combining large language models, voice services, and an interactive interface to support the technical assistance process within the Kabré HPC environment. The design is motivated by the heterogeneity of the user community, which implies that interactions cannot follow a single uniform behavior. Instead, the agent is expected to adapt its responses according to predefined user clusters such as social sciences, natural sciences, or computational disciplines, an approach consistent with recent work on LLM-based support systems in HPC environments [21] and findings in conversational agent research highlighting the importance of user characteristics and adaptive interaction design in socio-technical systems [6]. This adaptation affects both the level of technical detail and the explanatory style used in responses.

At the core of the system, the Gemma 4 E4B model is used, deployed within the Kabré infrastructure. This model is selected due to its ability to operate efficiently in high-performance computing environments and its strong performance in natural language understanding and generation tasks. These characteristics make it suitable for handling a wide range of technical support queries, from

basic system usage to more complex issues such as job execution, software installation, and resource management.

Interaction with the system is based exclusively on textual input from the user. Users communicate with the agent only through text, without any speech input capability. However, the system incorporates VoxCPM2 as a text-to-speech (TTS) engine for the agent’s output. In this configuration, the language model generates textual responses that are subsequently converted into speech, enabling a multimodal output experience while maintaining a simple and consistent text-only input interface. The use of speech synthesis in this context is motivated by advances in conversational neural TTS systems, which show that incorporating conversational context and semantic awareness significantly improves prosody naturalness and perceived realism in voice agents [11]. Such approaches demonstrate that high-quality conversational TTS benefits from modeling dialogue context and utterance-level dependencies, leading to more natural and contextually appropriate speech generation.

The user interface is developed using Unity, enabling a real-time interactive environment for communication with the agent. This interface serves as the main point of interaction between users and the support system, complementing the traditional ticket-based workflow. The use of an avatar-based representation of the agent is intended to make the interaction more intuitive and accessible, particularly for less technical users. This design choice is consistent with research on interactive agent-based avatar systems, where embodied virtual agents have been shown to improve usability and engagement by combining visual representation with structured behavioral models [4].

Additionally, the system may incorporate facial animation technologies such as the NVIDIA LipSync model to improve synchronization between generated speech and visual representation. This component is considered optional and is intended to enhance the naturalness of the interaction experience.

The system is designed to run primarily within the Kabré infrastructure, leveraging its computational resources for both the language model and back-end services. The IVA will be accessible exclusively through the official Kabré website, operating as an online service integrated into the institutional portal, where users can interact with the agent directly from their web browser.

Overall, this architecture is intended to address the main limitations identified in the current support system, including ticket overload, user heterogeneity, and long response times. Rather than replacing human support, the system is designed to complement it by automating repetitive queries while escalating complex cases to human experts when necessary.

7 Methodology

7.1 Evaluation Protocol Design and a Consistency Matrix

1. Definition of Conditions.

Experimental Condition:

The experimental condition introduces participants to the Intelligent Virtual Agent integrated into the Kabré supercomputer web portal. Participants interact with the agent through a browser-based interface featuring a 3D animated avatar developed in Unity, which provides visual expressions and synchronized speech to support communication.

The system uses an intelligent routing mechanism that operates entirely in the background: when a user submits a query, the backend automatically identifies its nature and context, such as routine software installation requests or complex job management issues, and routes it to the most appropriate processing component.

The IVA is designed to assist users in resolving common technical support scenarios within the Kabré infrastructure. Users interact with the system by submitting natural language queries in plain text, and the agent responds using a combination of structured textual explanations and voice synthesis powered by the *VoxCPM2* engine.

During the experiment, participants complete a set of pre-defined and standardized technical tasks representative of real support cases in the Kabré environment, including job submission issues (Slurm), software environment configuration using modules, and interpretation of system error logs. All participants perform the tasks concurrently under the same experimental session to ensure consistent exposure conditions.

The system operates in real time and is evaluated under controlled laboratory conditions, ensuring that all participants are exposed to the same set of tasks and interaction flow. No human support staff intervenes during the interaction, allowing the assessment of the IVA as a fully automated support interface.

Baseline Definition (Historical Ticketing Experience):

The baseline condition in this study is not a concurrently administered experimental group, but rather a retrospective and experiential baseline derived from participants' prior interactions with the Kabré technical support ticketing system. All participants are active users of the supercomputing infrastructure who have previously submitted and managed support tickets through the CeNAT helpdesk platform. As a result, they possess an internalized reference of system performance, including perceived waiting times, response delays, and interaction friction associated with the traditional asynchronous support workflow.

This study therefore adopts a within-subject design in which participants are exposed only to the Intelligent Virtual Agent (IVA), and all comparative judgments regarding efficiency, usability, and frustration are anchored against their prior real-world experience with the ticketing system.

Justification:

The selection of this baseline is fundamental to isolate and quantify the empirical impact of the Intelligent Virtual Agent's architecture and features on system performance and user experience. This study does not use a concur-

rent control group, but instead adopts a within-subject design based on a retrospective baseline derived from participants' prior experience with the Kabré technical support ticketing system.

This approach enables comparison between the interaction with the Intelligent Virtual Agent and the users' previously internalized experience, including perceived waiting times, response delays, and interaction friction accumulated through real-world use of the traditional support system.

Methodologically, the use of standardized technical tasks ensures controlled interaction with the automated system, such that response time, success rate, and interaction efficiency metrics reflect solely the performance of the agent. This eliminates the influence of delays inherent to the human-operated asynchronous support system and allows a more direct evaluation of the automated interface.

Furthermore, this design enables the assessment of whether the visual and auditory components of the interface contribute to reduced frustration and improved perceived usability when contrasted with users' prior experience with the traditional ticketing system. Consequently, the study provides evidence regarding the potential of the Intelligent Virtual Agent to improve efficiency and user experience in high-performance computing support environments.

2. Methodological Consistency Matrix.

The experimental design and its operational alignment are structured in Table 2

3. Adaptation of the Researcher's Protocol.

The experimental protocol was developed and implemented in HTML, integrating study identification, informed consent, randomized group assignment, task instructions, post-interaction questionnaires, and data anonymization policies. All materials can be accessed in the [project repository](#).

4. Theoretical Justification in Human-Computer Interaction (HCI).

The architectural and interaction design choices implemented in the Intelligent Virtual Agent for the Kabré supercomputer are grounded in foundational frameworks of HCI and social psychology, aligned with the methodological framework of the study.

The design of this virtual agent was not conceived as a mere visual enhancement, but as a solution to real interaction problems within the computing cluster. According to [5], the use of a stylized humanoid avatar instead of a plain text-based chat interface fulfills a key functional role. The moderate facial expressions and lip-syncing achieved through the VoxCPM2 engine help guide the conversation in a more natural way. By also communicating through non-verbal cues, the interface regulates dialogue pacing and makes support instructions easier to process, reducing the user's cognitive load and

Research Questions	Variable or Construct	Validated Instrument	Associated Task
RQ1: <i>To what extent can an Intelligent Virtual Agent improve the efficiency of technical support in the Kabré supercomputer environment, specifically minimizing user-side waiting latencies and optimizing request processing compared to traditional ticketing?</i>	Perceived Efficiency, Agent Competence, and Support Resolution Time	ARQx and objective metrics of resolution time	The participant interacts with the IVA to resolve a Slurm-related support scenario reflecting common historical ticketing cases in the Kabré environment, such as correcting an incorrect job submission parameter or checking the execution status of a job in the queue.
RQ2: <i>What is the perceived usability of an automated IVA across a user base with diverse levels of technical expertise, and how does the agent's operational autonomy influence user trust when navigating complex supercomputing tasks?</i>	System Usability, Trust, and Expertise-Based Experience	System Usability Scale (SUS) and qualitative user feedback	The participant uses the IVA to locate, load, and configure a scientific software dependency using environment modules, following the agent's guidance with minimal external assistance in order to evaluate perceived usability, trust, and system autonomy.
RQ3: <i>Does interacting with a 3D animated avatar that speaks and shows expressions actually help reduce user frustration, or does it feel like an unnecessary distraction compared to a simple, text-only chat?</i>	Perceived Frustration, Mental Demand, and Interface Animacy	NASA-TLX (Frustration and Mental Demand subscales) and Godspeed Questionnaire	The participant is presented with a complex and ambiguous system error log related to disk space or file access restrictions. The agent assists in interpreting the issue, enabling assessment of whether animated and vocal feedback reduces perceived frustration and cognitive load compared to raw textual logs.

avoiding discomfort associated with the *uncanny valley* thanks to a balanced visual design.

This visual support is directly complemented by the management of the user's emotional state. Based on the strategies described by [1], incorporating empathetic responses and a more natural conversational style is essential to reducing frustration in a highly complex environment such as supercomputing. The main advantage here is immediacy: unlike traditional ticketing systems, where a researcher may wait hours or days for a technical response in full uncertainty, the agent acts as a real-time companion precisely when the issue occurs. When code fails to compile or storage becomes blocked, the

agent responds instantly, acknowledges the user’s stress, provides an immediate solution, builds trust, and prevents task abandonment.

Finally, this trust does not depend only on what the agent says or when it says it, but also on the authority projected through its visual presence. Although the Proteus Effect analyzed by [yee] explains how a person’s behavior changes depending on their own avatar representation, its underlying psychological principles demonstrate that the visual design of a character strongly shapes the expectations of the observer. By unifying the agent’s identity with the cultural symbols of CeNAT’s computing nodes (Nukwa, Kurá, and Dribe), the interface naturally conveys a sense of institutional competence and technical expertise. This directly influences the researcher’s confidence, increasing trust in the system and improving perceived usability when configuring complex scientific variables and dependencies.

7.2 Exploratory Data Analysis

The data extracted from the Zammad SQL database of the Kabré supercomputer ticketing system are analyzed in order to identify frequent user queries and classify request types. The objective of this analysis is to derive meaningful support categories that enable the assignment of specialized agents to each type of request.

References

- [1] Timothy W Bickmore and Rosalind W Picard. “Establishing and maintaining long-term human-computer relationships”. In: *ACM Transactions on Computer-Human Interaction (TOCHI)* 12.2 (2005), pp. 293–327.
- [2] Brandon S Biggs and Kaylee Dalton. *Outcomes of HPC User Support using a Science Gateway AI Assistant*. Tech. rep. Idaho National Laboratory (INL), Idaho Falls, ID (United States), 2024.
- [3] Akhilesh Bondapalli et al. “AskHPC: A ChatBot for High Performance Computing User Support”. In: *Proceedings of the SC’25 Workshops of the International Conference for High Performance Computing, Networking, Storage and Analysis*. 2025, pp. 727–739.
- [4] Hongming Cai and Yuanjun He. “Interactive agent based user avatar in the virtual environment”. In: *Proceedings of the 2004 ACM SIGGRAPH international conference on Virtual Reality continuum and its applications in industry*. 2004, pp. 100–103.
- [5] Justine Cassell et al. “Designing embodied conversational agents”. In: *Embodied conversational agents* (2000), pp. 29–63.
- [6] Stephan Diederich et al. “On the design of and interaction with conversational agents: An organizing and assessing review of human-computer interaction research”. In: *Journal of the Association for Information Systems* 23.1 (2022), pp. 96–138.
- [7] Ali Doosthosseini et al. “SAIA: A seamless slurm-native solution for HPC-based services”. In: *The Journal of Supercomputing* 82.7 (2026), p. 403.

- [8] Diego Omar Encinas et al. “An agent-based model for analyzing the HPC input/output system”. In: (2020).
- [9] Chris Faigle et al. “Integrating virtual environments with high performance computing”. In: *Proceedings of IEEE Virtual Reality Annual International Symposium*. IEEE. 1993, pp. 62–68.
- [10] Yuping Fan et al. “Deep reinforcement agent for scheduling in hpc”. In: *2021 IEEE International Parallel and Distributed Processing Symposium (IPDPS)*. IEEE. 2021, pp. 807–816.
- [11] Haohan Guo et al. “Conversational end-to-end tts for voice agents”. In: *2021 IEEE Spoken Language Technology Workshop (SLT)*. IEEE. 2021, pp. 403–409.
- [12] Shubham Jain, Amit Gupta, and Kumari Neha. “AI enhanced ticket management system for optimized support”. In: *Proceedings of the 4th International Conference on AI-ML Systems*. 2024, pp. 1–7.
- [13] Aparajita Jaiswal et al. “Empa: An AI-Powered Virtual Mentor for Developing Global Collaboration Skills in HPC Education”. In: *arXiv preprint arXiv:2511.17669* (2025).
- [14] Juan Monroy et al. “Improving performance in HPC centers by modeling users through an evolutionary virtual interface”. In: *IEEE Transactions on Instrumentation and Measurement* 57.9 (2008), pp. 1885–1893.
- [15] Juan Monroy et al. “Intelligent Virtual Interface for Improving Performance in HPC Centers by Modelling Users and their Satisfaction”. In: *2006 IEEE Symposium on Virtual Environments, Human-Computer Interfaces and Measurement Systems*. IEEE. 2006, pp. 69–74.
- [16] Juan Monroy et al. “Intelligent Virtual Interface for Improving Performance in HPC Centers by Modelling Users and their Satisfaction”. In: *2006 IEEE Symposium on Virtual Environments, Human-Computer Interfaces and Measurement Systems*. IEEE. 2006, pp. 69–74.
- [17] Jargalsaikhan Narantuya et al. “Multi-agent deep reinforcement learning-based resource allocation in hpc/ai converged cluster”. In: (2022).
- [18] Gennady A Oparin, Vera Bogdanova, and Anton A Pashinin. “Automated tools for the development of microservice compositions for hybrid scientific computations.” In: *ICCS-DE*. 2020, pp. 201–213.
- [19] Daniel Rosendo et al. “AI Agents for Enabling Autonomous Experiments at ORNL’s HPC and Manufacturing User Facilities”. In: *Proceedings of the SC’25 Workshops of the International Conference for High Performance Computing, Networking, Storage and Analysis*. 2025, pp. 2354–2361.
- [20] WGAGP Sanjeewa. “IT Ticketing System with a Chatbot”. PhD thesis. 2021.
- [21] Jochen Wulf and Juerg Meierhofer. “Exploring the potential of large language models for automation in technical customer service”. In: *arXiv preprint arXiv:2405.09161* (2024).
- [22] Junqi Yin et al. “chatHPC: Empowering HPC users with large language models: J. Yin et al.” In: *The Journal of Supercomputing* 81.1 (2025), p. 194.

- [23] Vladimir S Zaborovsky et al. “Improving Efficiency of Hybrid HPC Systems Using a Multi-agent Scheduler and Machine Learning Methods”. In: *Supercomputing Frontiers and Innovations* 10.2 (2023), pp. 104–126.